

Comparison of Eye-Tracking Representations for Engagement Inference in Office Workers

[To COMPLETE: Author name]

[To COMPLETE: Institution]

[To COMPLETE: City], [To COMPLETE: Country]

author@example.org

Abstract

Keywords

eye tracking, engagement inference, representation learning, GazeMAE, MOMENT, subject generalisation

1 Introduction

- Engagement inference in office work
- Eye tracking: gaze behaviour, pupil dynamics, behavioural evidence
- Practical representation choice for eye-tracking windows
- Raw temporal signals: information retention, high dimensionality
- Handcrafted features: domain knowledge, interpretability
- Learned embeddings: transfer promise, preprocessing and deployment cost
- Gaze-specific GazeMAE; general time-series MOMENT
- Evidence gap: like-for-like comparison on identical windows and labels
- Unseen-person generalisation; subject-disjoint evaluation
- Research question: representation reliability for self-reported engagement in unseen office workers
- Controlled comparison: raw, handcrafted, GazeMAE, MOMENT
- Frozen pretrained encoders; trained prediction head only
- Contribution: leakage-safe LOSO evaluation
- Contribution: practical representation-selection guidance

2 Related Work

- Eye tracking for engagement, cognitive load, affect, and user-state inference
- Gaze and pupil feature engineering
- Gaze representation learning; GazeMAE [1]
- General-purpose time-series foundation models; MOMENT [2]
- Frozen encoder transfer; out-of-the-box downstream utility
- Subject effects and overoptimistic random window splits
- Subject-disjoint evaluation and fold-local preprocessing
- Study position: empirical representation comparison; no new encoder or fine-tuning method

3 Methodology

3.1 Data and Eye-Tracking Representations

- TrustMe study; office-worker recordings

- Participant count, sessions, recording duration, sampling rate, exclusions
- Tobii channels: left/right pupil size, gaze x/y , average distance
- q5 engagement item: “How immersed in your work did you feel just now?”
- q5 response scale and observed raw-label distribution
- Window duration, overlap, temporal ordering
- Common window_uid intersection across all four representations
- Labelled windows only; final subject and window counts
- Validity fields; missing-value and quality-control policy
- Display-resolution metadata; per-recording x/y normalization to $[0, 1]$
- Out-of-screen coordinates: retained for modelling policy; excluded from descriptive heatmap
- Representation order: raw, handcrafted features, GazeMAE embeddings, MOMENT embeddings
- Raw: five channels, 120 samples per channel, truncate/pad policy, final dimensionality
- Handcrafted: feature families, metadata removal, missingness filtering, correlation filtering
- GazeMAE: checkpoint, inputs, embedding selection, dimensionality, frozen encoder
- MOMENT: checkpoint, inputs, pooling, dimensionality, frozen encoder
- **Figure 1 — normalized gaze-density map:** screen-coordinate coverage; recording-quality context; non-predictive descriptive evidence

3.2 Experimental Design

- Primary task: binary q5 engagement inference
- LOSO target construction: training-subject mean centring; held-out subject global-training-mean centring
- Diagnostic visualization label only: raw q5; binary q5 above global processed-data median; no use in model training or metrics
- Same aligned cohort, target definition, and folds across representations
- Majority-class baseline
- Random Forest: 300 trees, balanced class weights, fixed seed
- MLP: GELU, layer normalization, Adam, early stopping, fixed seed
- Representation-by-model grid; robustness check rather than model-comparison claim
- Fold-local handcrafted-feature selection and imputation
- Fold-local median imputation; train-fit standardization; held-out-subject transform
- No test-subject information in preprocessing parameters
- Accuracy: primary metric

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

Information Society 2026, Ljubljana, Slovenia

© 2026 Copyright held by the owner/author(s).

- Accuracy delta and balanced-accuracy delta: paired majority-baseline comparison within LOSO folds
- Macro-F1 and ROC-AUC: secondary discrimination metrics
- Subject-macro aggregation
- One-class test folds: retained accuracy; undefined balanced accuracy, macro-F1, and AUC; reported valid-fold counts
- Persistence baseline: contextual within-subject temporal check; text only; not comparable with unseen-subject main table

4 Results

- **Table 1 — primary LOSO results:** majority baseline; raw, handcrafted, GazeMAE, MOMENT; Random Forest and MLP
- Table columns: representation; model; accuracy (Δ majority); balanced accuracy (Δ majority); macro-F1; ROC-AUC
- Table order: majority baseline; raw; handcrafted; GazeMAE; MOMENT
- Subject-macro estimates; valid-fold counts in caption or table note
- Bold: best value per metric; no overemphasis on small model differences
- Main result: highest accuracy; absolute and majority-baseline improvement
- Representation result: raw and handcrafted performance relative to frozen GazeMAE and MOMENT embeddings
- Model result: consistency of representation ranking across Random Forest and MLP
- One-class held-out-subject treatment; metric availability
- Persistence-check result: contextual temporal predictability; separate protocol; no direct ranking against LOSO models
- **Figure 2 — selected projection:** one representation-method combination; raw q5 continuous colour scale
- Figure-selection criterion: clearest engagement gradient or clearest absence of label separation
- Figure interpretation: exploratory geometry; mixed colours as limited label separability; no causal or performance claim

5 Discussion

- Direct answer to the research question
- Frozen learned embeddings: out-of-the-box utility relative to raw and handcrafted alternatives
- Representation effect; secondary role of classifier choice
- Practical recommendation: accuracy, preprocessing burden, storage, computation, interpretability
- Implication of non-improving embeddings: task–pretraining mismatch; need for domain adaptation or fine-tuning
- Projection and heatmap: descriptive support only; no separability proof
- Persistence result: temporal autocorrelation; distinct from cross-person generalisation
- Limitations: single dataset, office-worker sample, self-report noise, binary target construction
- Limitations: frozen checkpoints, fixed windowing, quality-control choices, limited subject count
- Future work: multi-dataset validation, calibrated person-alisation, adaptation/fine-tuning, alternative target formulations

6 Conclusion

- Controlled four-representation comparison for unseen-worker engagement inference
- Final best-performing representation and headline accuracy
- Frozen-encoder conclusion; out-of-the-box deployment recommendation
- Concise practitioner guidance: representation choice under limited labelled data

Acknowledgements

References

- [1] Louise Gillian C. Bautista and Jr. Naval Prospero C. 2020. GazeMAE: general representations of eye movements using a micro-macro autoencoder. *arXiv preprint arXiv:2009.02437*. <https://arxiv.org/abs/2009.02437>.
- [2] Mononito Goswami, Konrad Szafer, Arjun Choudhry, Yifu Cai, Shuo Li, and Artur Dubrawski. 2024. MOMENT: a family of open time-series foundation models. *arXiv preprint arXiv:2402.03885*. <https://arxiv.org/abs/2402.03885>.